

Interim Report: Theoretical and Empirical Analysis of Gradient Centralization

Kabil

Department of Computer Science
Jadavpur University
Kolkata, India

Teammate Name 2

Department of Computer Science
Jadavpur University
Kolkata, India

Teammate Name 3

Department of Computer Science
Jadavpur University
Kolkata, India

Abstract—This report presents an interim analysis of Gradient Centralization (GC), a simple yet effective optimization technique for Deep Neural Networks. We provide a rigorous theoretical formulation demonstrating that GC acts as an implicit diagonal-free preconditioner, explicitly reducing the condition number of the effective Hessian. Furthermore, we provide empirical validation on the CIFAR-10 dataset, demonstrating GC’s superior convergence properties compared to standard Adam.

I. INTRODUCTION

Optimization in Deep Learning often suffers from poor conditioning of the loss landscape. While adaptive methods like Adam scale gradients using moving averages of past moments, they do not explicitly target the condition number associated with the scale invariance of network weights. Gradient Centralization (GC) operates by centralizing the gradient vectors to have zero mean. In this report, we formally prove why GC accelerates convergence by analyzing its effect on the Hessian matrix.

II. THEORETICAL ANALYSIS

A. Gradient Centralization as a Projection Operator

Let $\mathcal{L}(\mathbf{w})$ be the loss function parameterized by weights $\mathbf{w} \in \mathbb{R}^N$. The GC operation Φ_{GC} on the gradient $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$ is defined as:

$$\Phi_{GC}(\mathbf{g}) = \mathbf{g} - \mu_{\mathbf{g}} \mathbf{1} \quad (1)$$

where $\mu_{\mathbf{g}} = \frac{1}{N} \sum_{i=1}^N g_i$ is the mean of the gradient and $\mathbf{1} \in \mathbb{R}^N$ is a vector of ones. This operation can be written using a projection matrix P :

$$P = I - \frac{1}{N} \mathbf{1} \mathbf{1}^T \quad (2)$$

P is an orthogonal projection matrix ($P^T = P$ and $P^2 = P$). Thus, $\Phi_{GC}(\mathbf{g}) = P \nabla \mathcal{L}(\mathbf{w})$.

B. Effective Hessian and Condition Number Reduction

The condition number of the Hessian matrix $H = \nabla^2 \mathcal{L}(\mathbf{w})$ governs the convergence rate. With GC, the effective geometry of the loss landscape is governed by the projected Hessian:

$$\tilde{H} = PHP \quad (3)$$

For networks with homogeneous activations, the weight scale introduces a massive eigenvalue in the direction of the mean

vector $\mathbf{1}$. Because P projects out the subspace spanned by $\mathbf{1}$, applying P eliminates this extreme curvature:

$$\lambda_{\max}(\tilde{H}) \leq \lambda_{\max}(H) \quad \text{and} \quad \lambda_{\min}^+(\tilde{H}) \geq \lambda_{\min}(H) \quad (4)$$

Therefore, the effective condition number $\kappa(\tilde{H})$ is strictly improved:

$$\kappa(\tilde{H}) = \frac{\lambda_{\max}(\tilde{H})}{\lambda_{\min}^+(\tilde{H})} < \frac{\lambda_{\max}(H)}{\lambda_{\min}(H)} = \kappa(H) \quad (5)$$

This formally demonstrates GC’s role as a preconditioner without explicit $O(N^3)$ matrix inversion.

C. Linear Convergence Rates for GC

Assume $\mathcal{L}(\mathbf{w})$ is μ -strongly convex and L -smooth. The effective smoothness becomes $\tilde{L} = \lambda_{\max}(PHP) \leq L$, and effective strong convexity $\tilde{\mu} = \lambda_{\min}^+(PHP) \geq \mu$. By setting the learning rate $\alpha = \frac{2}{\tilde{L} + \tilde{\mu}}$, we obtain a linear convergence rate:

$$\|\mathbf{w}_{t+1} - \mathbf{w}^*\| \leq \left(\frac{\kappa(\tilde{H}) - 1}{\kappa(\tilde{H}) + 1} \right) \|\mathbf{w}_t - \mathbf{w}^*\| \quad (6)$$

Since $\kappa(\tilde{H}) < \kappa(H)$, the contraction factor is strictly smaller, proving that GC accelerates convergence.

D. Extension to Non-Convex Settings

For non-convex settings with $\tilde{L}$ -Lipschitz continuous projected gradients, GC guarantees convergence to a first-order stationary point at a rate of $O(1/T)$. Furthermore, because the dominant positive curvature is removed, GC facilitates faster escape from saddle points, converging to second-order stationary points more effectively.

III. EXPERIMENTAL RESULTS

To validate our theoretical findings, we trained a Convolutional Neural Network (CNN) on the CIFAR-10 dataset. We compared the standard Adam optimizer against our custom implementation of Adam with Gradient Centralization (AdamGC).

During training, we tracked the Maximum Eigenvalue ($\lambda_{\max}$) of the Hessian using the Power Iteration method with Hessian-Vector Products (HVPs).

Figure 1 illustrates our updated output results. Figure 2 provides a side-by-side comparison. As theoretically predicted,



Fig. 1. New Output Results demonstrating the performance metrics of the CNN model.




Fig. 2. Comparative analysis between Standard Adam and AdamGC. The plots show Training Loss, Training Accuracy, and the empirical Maximum Eigenvalue (λ_{max}) of the Hessian across 10 epochs.

AdamGC successfully suppresses the extreme curvature (lower λ_{max}) compared to standard Adam, resulting in a smoother loss landscape and accelerated convergence.

IV. CONCLUSION

We have formally proven that Gradient Centralization acts as a computationally efficient preconditioner. By reducing the condition number of the Hessian matrix, GC accelerates optimization in both strongly convex and non-convex environments. Our empirical results on CIFAR-10 strongly corroborate the theoretical derivations.